

1. **Bias and fairness:** Write a bias and fairness report covering, at a minimum:

- Any known sampling or coverage limitations in your data.
- Any proxy labels you used and the risks associated with them.
- How errors might be distributed unevenly across different conditions or time periods.

- Our traffic volume dataset only represents one single highway corridor over a specific historical timeline. It does not capture data from smaller residential estate roads, side streets, or new shortcut lanes. Because the data lacks of local map coverage, an AI system trained on this data will be highly biased toward highway travel patterns and will fail to make accurate traffic recommendations for everyday city street driving.

2. **Governance and sustainability:** Include a short discussion of governance and sustainability in mobility AI. Consider:

- What oversight should be in place before a model is trusted for real-world decision-making.
- The environmental or resource trade-offs associated with your modelling choices.

- Before this traffic AI model can be trusted to run live on city dashboards, We cannot let an automated script change public highway electronic signs by itself. We need an independent panel of traffic engineers to run simulation stress tests on the code to verify the accuracy limits. There must also be a clear emergency kill switch so control operators can shut down the AI recommendations instantly if the system starts throwing erratic alerts during a real crisis.